

Hello,

The task at hand is a binary classification problem, for which we provide both a training and a validation set as csv files.

We would like you to work on this classification task by building a classifier on the training data, maximising performance on the validation data, the target variable is 'y'. Although we work with Python, we are fine with you solving the task with R.

We expect to receive the source code of the solution and a write up with your observations, thought process, any interesting findings and a description of the results presented.

We hope you will have fun!

Good Luck,  
The Grover Data Team